

ECON7720Project_ButtsCullenDouglasKarnatzLadina

December 2, 2024

ECON7720 Final Project *Nick Butts, Ryan Cullen, Judd Douglas, Matt Karnatz, Collin Ladina*

1 Dataset Overview and Project Description

1.0.1 This problem involves the Carseats data set data set containing sales of child car seats at 400 different stores. A number of characteristics are recorded:

- Sales: Unit sales (in thousands) at each location
- CompPrice: Price charged by competitor at each location
- Income: Community income level (in thousands of dollars)
- Advertising: Local advertising budget for company at each location (in thousands of dollars)
- Population: Population size in region (in thousands)
- Price: Price company charges for car seats at each site
- ShelfLoc: A factor with levels Bad, Good and Medium indicating the quality of the shelving location for the car seats at each site
- Age: Average age of the local population
- Education: Education level at each location
- Urban: A factor with levels No and Yes to indicate whether the store is in an urban or rural location
- US: A factor with levels No and Yes to indicate whether the store is in the US or not

1.0.2 Please build a model to understand the determination of sales using the first 300 observations.

2 Libraries, Ignore Some Warnings, and Data

```
[11]: import numpy as np
import pandas as pd
from ISLP import load_data

import warnings
warnings.filterwarnings("ignore")

Carseats = load_data('Carseats')
data = Carseats.head(300)

data.info() # data.shape[0] for number of rows, data.shape[1] for number of
↪ columns
```

```
numeric_data = data.select_dtypes(include=['int64','float64'])
categorical_data = data.select_dtypes(include='category')
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 300 entries, 0 to 299
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sales           300 non-null   float64
1   CompPrice       300 non-null   int64
2   Income         300 non-null   int64
3   Advertising     300 non-null   int64
4   Population      300 non-null   int64
5   Price          300 non-null   int64
6   ShelfLoc       300 non-null   category
7   Age            300 non-null   int64
8   Education       300 non-null   int64
9   Urban          300 non-null   category
10  US             300 non-null   category
dtypes: category(3), float64(1), int64(7)
memory usage: 20.1 KB
```

```
[2]: print("Duplicates in numeric data:", numeric_data.duplicated().sum())
      print("\nNumber of nulls for each column:\n",data.isnull().sum())
      print("\nNumber of infinite values for each column:\n\n",np.isinf(numeric_data).
      ↪sum())
```

Duplicates in numeric data: 0

Number of nulls for each column:

```
Sales           0
CompPrice       0
Income         0
Advertising     0
Population      0
Price          0
ShelfLoc       0
Age            0
Education       0
Urban          0
US             0
dtype: int64
```

Number of infinite values for each column:

```
Sales           0
CompPrice       0
Income         0
```

```
Advertising    0
Population     0
Price          0
Age            0
Education      0
dtype: int64
```

From the 300 observations we are using as our dataset, we see that we have no missing values, infinite values or duplicated rows in our dataset. There are 7 numerical variables and 3 category variables being ShelfLoc, Urban, and US.

```
[3]: data.describe().transpose()
```

```
[3]:
```

	count	mean	std	min	25%	50%	75%	\
Sales	300.0	7.362967	2.835229	0.0	5.255	7.445	9.11	
CompPrice	300.0	124.536667	15.343840	77.0	115.000	123.500	135.00	
Income	300.0	69.843333	27.716291	21.0	44.750	70.500	91.25	
Advertising	300.0	6.036667	6.262538	0.0	0.000	5.000	11.00	
Population	300.0	260.716667	147.445959	10.0	136.250	265.500	391.50	
Price	300.0	115.490000	23.655199	24.0	99.750	117.000	131.00	
Age	300.0	53.743333	16.298761	25.0	39.750	55.500	66.00	
Education	300.0	14.020000	2.702581	10.0	12.000	14.000	16.25	


	max
Sales	14.9
CompPrice	162.0
Income	120.0
Advertising	25.0
Population	509.0
Price	191.0
Age	80.0
Education	18.0

3 Visualizing Target Variable (Sales)

```
[4]: from matplotlib import pyplot as plt
import seaborn as sns


sns.histplot(data['Sales'],color='blue')
plt.title(f"{'Sales'}")
plt.xlabel("Value")
plt.ylabel("Frequency")
```

```
[4]: Text(0, 0.5, 'Frequency')
```



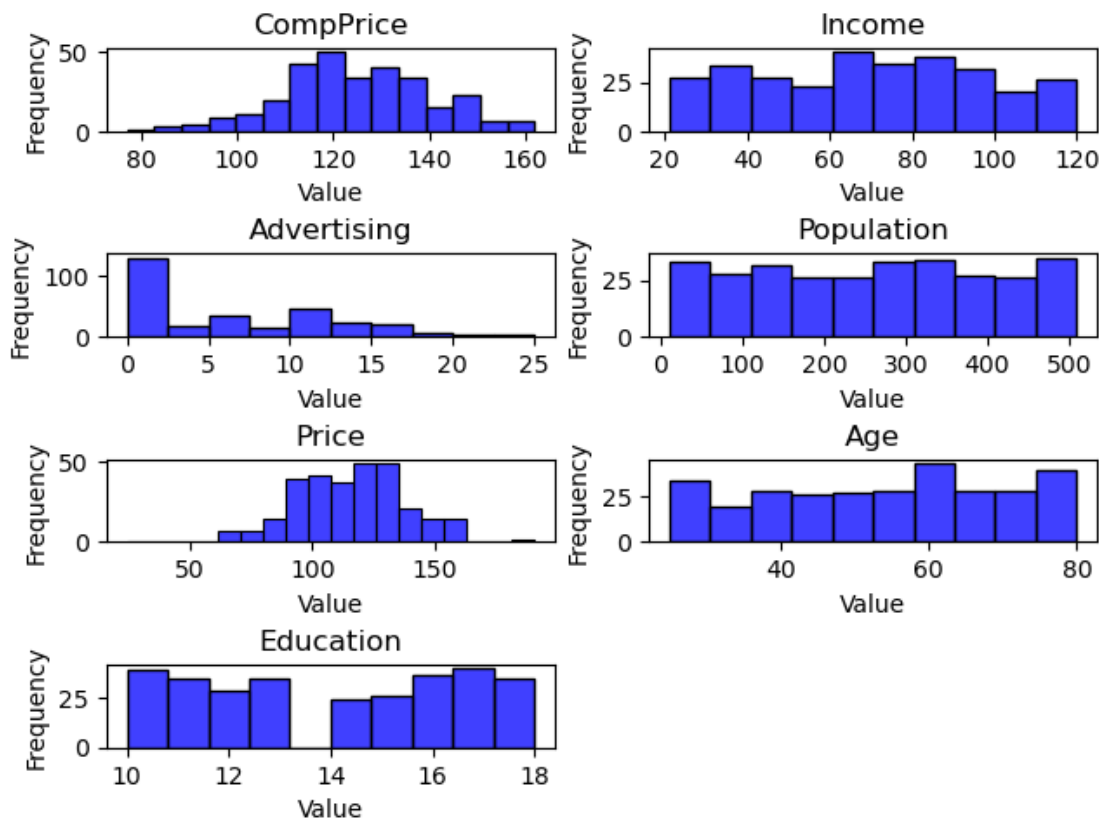
4 Visualizing Categorical Variables

```
[5]: plt.figure(figsize=(10,10))
for i, feature in enumerate(categorical_data.columns):
    counts = data[feature].value_counts().head(10)
    plt.subplot(2, 2, i+1)
    sns.barplot(x=counts.values, y=counts.index, palette='rocket_r')
    plt.title(f"{feature}")
    plt.xlabel("Count")
plt.tight_layout(pad=1.0)
#plt.savefig("categorical_features_plots.png"); plt.clf();
```



5 Visualizing Numeric Features/Inputs

```
[6]: for i, feature in enumerate(numeric_data.columns[1:]): # Sales is first column,
      ↪ so it is excluded since it is target variable
      plt.subplot(4,2,i+1) # rows (first param) * cols (second param) = 8
      sns.histplot(data[feature],color='blue')
      plt.title(f"{feature}")
      plt.xlabel("Value")
      plt.ylabel("Frequency")
plt.tight_layout(pad=0.5)
#plt.savefig("continuous_features_plots.png"); plt.clf()
```



From some summary statistics and exploratory data analysis, the average Sales from each location is 7.326 (in thousands). From the bell-shaped curve of our target variable we see that the distribution of sales is approximately normal. The majority of Store Locations (observations) from our dataset show that the shelving location was medium, in a urban location and stores were located in the US. From visualizing the numerical features distributions apart of the dataset, CompPrice and Price can be assumed to be approximately normal. Advertising is skewed right, with education, income, population, and age all being uniformly distributed.

6 Data Pre-Processing

1. Scale numeric features with RobustScaler (robust to outliers by scaling with 25th-75th percentiles of dataset)
2. Encode categorical variables with OneHotEncoder
 - 1 and 2 are combined in ColumnTransformer (alternative could be using Pipeline object storing all preprocessing)
3. Split into training and testing

```
[7]: from sklearn.compose import ColumnTransformer
      from sklearn.model_selection import train_test_split
```

```

from sklearn.preprocessing import OneHotEncoder, RobustScaler
from sklearn.pipeline import Pipeline

X = data.drop('Sales',axis=1)
y = data['Sales']

numeric_features = ['CompPrice', 'Income', 'Advertising', 'Population','Price',
    ↪ 'Age', 'Education']
categorical_features = ['ShelveLoc', 'Urban', 'US']

preprocessor = ColumnTransformer(transformers = [('num', RobustScaler(),
    ↪ numeric_features),
    ('cat', OneHotEncoder(sparse_output=False), categorical_features) ])

X = preprocessor.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

```

7 Models (Uses K-Fold Cross-Validation With $K = 5$)

7.0.1 1. Train model with GridSearchCV

7.0.2 2. Predict testing data with best model from step 1

7.0.3 3. Metrics: R-squared for all, Correlation Matrix in Regression, Feature Importance in Ensembles

7.0.4 Linear Regression

```

[8]: import sklearn.model_selection as skm
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.linear_model import LinearRegression, Ridge, Lasso

kfold = skm.KFold(5, random_state=42, shuffle=True)

linear_grid = skm.GridSearchCV(
    LinearRegression(),
    {}, # no hyperparameter tuning
    refit=True,
    cv=kfold,
    scoring='accuracy')

linear_grid.fit(X_train,y_train)
best_linear = linear_grid.best_estimator_

y_pred_train = best_linear.predict(X_train)

```

```

y_pred_test = best_linear.predict(X_test)

coefficients_rounded = [f'{num:.4f}' for num in best_linear.coef_]

print(f"Intercept: {best_linear.intercept_:.4f}\nCoefficients:␣
↪{coefficients_rounded}")
print(f"Training R-Squared: {r2_score(y_train,y_pred_train):.4f}")

print(f'Testing R-Squared: {r2_score(y_test, y_pred_test):.4f}')

```

```

Intercept: 7.1272
Coefficients: ['1.9296', '0.7849', '1.6044', '0.0029', '-2.9583', '-1.1465',
'-0.1119', '-2.1860', '2.5370', '-0.3509', '-0.0544', '0.0544', '0.2406',
'-0.2406']
Training R-Squared: 0.8786
Testing R-Squared: 0.8887

```

From our linear regression model in predicting the sales of carseats, we see after cross validation our training set model has a R-squared value of 0.8786. In testing our model accuracy we see that are model performed very well. When tested against the best training model, we see that the testing R-squared value is 0.8887. We can conclude that 88.87% of the variation in Carseat sales can be explained by our independent variables.

7.0.5 Ridge Regression With Parameter Tuning

For Ridge and LASSO, alpha would not work in param_grid in GridSearchCV, so for loop was implemented

```

[9]: columns = ['Model', 'Intercept', 'Training R-Squared', 'Testing_
↪R-Squared', 'Coefficients']
ridge_models = pd.DataFrame(index=[0,1,2], columns=columns)
ridge_models.loc[0, 'Model'] = Ridge(alpha=5)
ridge_models.loc[1, 'Model'] = Ridge(alpha=1)
ridge_models.loc[2, 'Model'] = Ridge(alpha=0.1)

for i in ridge_models.index:
    ridge_grid = skm.GridSearchCV(
        ridge_models.loc[i, 'Model'],
        {}, # no hyperparameter tuning in grid search cv
        refit=True,
        cv=kfold,
        scoring='accuracy')

    ridge_grid.fit(X_train,y_train)
    best_ridge = ridge_grid.best_estimator_

    y_pred_train = best_ridge.predict(X_train)
    y_pred_test = best_ridge.predict(X_test)

```



```

ridge_models.loc[i, 'Intercept'] = f'{best_ridge.intercept_:.4f}'
ridge_models.loc[i, 'Coefficients'] = [f'{num:.4f}' for num in best_ridge.
↪coef_]
ridge_models.loc[i, 'Training R-Squared'] = ↵
↪f'{r2_score(y_train, y_pred_train):.4f}'
ridge_models.loc[i, 'Testing R-Squared'] = f'{r2_score(y_test, y_pred_test):.
↪4f}'

ridge_models

```

```

[9]:
      Model Intercept Training R-Squared Testing R-Squared \
0  Ridge(alpha=5)    7.1281            0.8725            0.8805
1  Ridge(alpha=1)    7.1272            0.8783            0.8881
2  Ridge(alpha=0.1)  7.1272            0.8786            0.8887

      Coefficients
0  [1.7514, 0.7374, 1.3991, -0.0079, -2.7286, -1...
1  [1.8911, 0.7748, 1.5573, 0.0008, -2.9091, -1.1...
2  [1.9257, 0.7839, 1.5995, 0.0027, -2.9533, -1.1...

```

Using Ridge Regression, we see that the model using an alpha set at 0.1 performed the best with a training and testing R-squared of 0.8786 and 0.8887 respectively.

7.0.6 LASSO Regression With Parameter Tuning

```

[10]: columns = ['Model', 'Intercept', 'Training R-Squared', 'Testing ↵
↪R-Squared', 'Coefficients']
lasso_models = pd.DataFrame(index=[0,1,2], columns=columns)
lasso_models.loc[0, 'Model'] = Lasso(alpha=0.1)
lasso_models.loc[1, 'Model'] = Lasso(alpha=0.01)
lasso_models.loc[2, 'Model'] = Lasso(alpha=0.0001)

for i in lasso_models.index:
    lasso_grid = skm.GridSearchCV(
        lasso_models.loc[i, 'Model'],
        {}, # no hyperparameter tuning in grid search cv
        refit=True,
        cv=kfold,
        scoring='accuracy')

    lasso_grid.fit(X_train, y_train)
    best_lasso = lasso_grid.best_estimator_

    y_pred_train = best_lasso.predict(X_train)
    y_pred_test = best_lasso.predict(X_test)

    lasso_models.loc[i, 'Intercept'] = f'{best_lasso.intercept_:.4f}'

```

```

lasso_models.loc[i, 'Coefficients'] = [f'{num:.4f}' for num in best_lasso.
↪coef_]
lasso_models.loc[i, 'Training R-Squared'] =
↪f'{r2_score(y_train, y_pred_train):.4f}'
lasso_models.loc[i, 'Testing R-Squared'] = f'{r2_score(y_test, y_pred_test):.
↪4f}'

lasso_models

```

```

[10]:
      Model Intercept Training R-Squared Testing R-Squared \
0  Lasso(alpha=0.1)      6.8200           0.8411           0.8582
1  Lasso(alpha=0.01)     6.6507           0.8779           0.8898
2  Lasso(alpha=0.0001)   6.6916           0.8786           0.8887

      Coefficients
0  [1.5302, 0.5154, 0.9997, 0.0000, -2.4998, -0.8...
1  [1.8878, 0.7572, 1.4796, 0.0000, -2.9160, -1.1...
2  [1.9292, 0.7846, 1.6031, 0.0027, -2.9579, -1.1...

```

Using Lasso Regression, we see that the model using an alpha set at 0.0001 performed the best with a training and testing R-squared of 0.8786 and 0.8887 respectively.


7.0.7 Correlation Matrix

```

[11]: encoded_categories = preprocessor.named_transformers_['cat'].
↪get_feature_names_out(categorical_features)
column_names = numeric_features + list(encoded_categories)
transformed_df = pd.DataFrame(X, columns=column_names)
transformed_df['Sales'] = y.reset_index(drop=True)

plt.figure(figsize=(12, 10))
sns.heatmap(transformed_df.corr(), annot=True, fmt=".2f", linewidths=0.5,
↪cmap="coolwarm")
plt.title("Correlation Matrix")
plt.show()

```



7.0.8 Random Forest With Parameter Tuning

```
[12]: from sklearn.ensemble import RandomForestRegressor

rf_grid = skm.GridSearchCV(
    RandomForestRegressor(random_state=42),
    param_grid={
        'ccp_alpha': [0.01, 0.1, 10],
        'max_features': [3, 5, None],
        'min_samples_split': [5, 10, 50]
    },
    refit=True,
    cv=kfold,
    scoring='accuracy')
```

```

rf_grid.fit(X_train,y_train)
best_rf = rf_grid.best_estimator_

y_pred_train = best_rf.predict(X_train)
y_pred_test = best_rf.predict(X_test)

print(f"Best parameters after tuning/cv: {rf_grid.best_params_}")
print(f"Training R-Squared: {r2_score(y_train,y_pred_train):.4f}")

print(f'Testing R-Squared: {r2_score(y_test, y_pred_test):.4f}')

feature_importances = best_rf.feature_importances_
features = transformed_df.drop('Sales',axis=1).columns

pd.DataFrame({
    'Feature': features,
    'Importance': feature_importances
}).sort_values(by='Importance', ascending=False)

```

```

Best parameters after tuning/cv: {'ccp_alpha': 0.01, 'max_features': 3,
'min_samples_split': 5}
Training R-Squared: 0.8751
Testing R-Squared: 0.6346

```

```

[12]:

```

	Feature	Importance
4	Price	0.205083
8	ShelveLoc_Good	0.173927
5	Age	0.095422
1	Income	0.092040
0	CompPrice	0.089748
2	Advertising	0.084315
7	ShelveLoc_Bad	0.075530
3	Population	0.069298
6	Education	0.041614
9	ShelveLoc_Medium	0.033297
12	US_No	0.011867
11	Urban_Yes	0.010020
10	Urban_No	0.009192
13	US_Yes	0.008646

From our random forest model setting alpha as 0.01 we see that our training R-squared value is high at 0.8751 but our testing R-squared value was 0.6346 which shows that 63.46% of our variability in Sales can be explained by our independent variables. Also from our random forest model we see that our top three features in our model were Price, Good Shelving Location, and Age.

7.0.9 XGBoost With Parameter Tuning

```
[13]: import xgboost as xgb

xgb_grid = skm.GridSearchCV(
    xgb.XGBRegressor(),
    param_grid={
        'max_features': [3,5,None],
        'min_samples_split': [5,10,50]
    },
    refit=True,
    cv=kfold,
    scoring='accuracy')

xgb_grid.fit(X_train,y_train)
best_xgb = xgb_grid.best_estimator_

y_pred_train = best_xgb.predict(X_train)
y_pred_test = best_xgb.predict(X_test)

print(f"Best parameters after tuning/cv: {xgb_grid.best_params_}")
print(f"Training R-Squared: {r2_score(y_train,y_pred_train):.4f}")

print(f'Testing R-Squared: {r2_score(y_test, y_pred_test):.4f}')

feature_importances = best_xgb.feature_importances_
features = transformed_df.drop('Sales',axis=1).columns

pd.DataFrame({
    'Feature': features,
    'Importance': feature_importances
}).sort_values(by='Importance', ascending=False)
```

Best parameters after tuning/cv: {'max_features': 3, 'min_samples_split': 5}

Training R-Squared: 1.0000

Testing R-Squared: 0.7026

```
[13]:
```

	Feature	Importance
8	ShelveLoc_Good	0.743838
7	ShelveLoc_Bad	0.114637
4	Price	0.056540
2	Advertising	0.021581
5	Age	0.021372
0	CompPrice	0.015734
1	Income	0.010088
6	Education	0.005808
10	Urban_No	0.004618
3	Population	0.003483

9	ShelveLoc_Medium	0.001313
12	US_No	0.000987
11	Urban_Yes	0.000000
13	US_Yes	0.000000

From our gradient boosted tree our testing R-squared value was 0.7026. We also see that the shelving location was the most important feature in our model.